Member-only story

Top 10 scenario-based interview questions with answer related to AWS Part-5

Deep · Follow
6 min read · Dec 28, 2023

It took me a lot of time and effort to provide these detailed answers, and Medium doesn't pay for technical articles like this. If you find these notes helpful, **please support me**!

👉 Click here to **buy me a coffee**, and don't forget to **follow me on Medium!**

On a free Medium plan? **Click here & Please Do Follow Me!**

```
I would really appreciate it
if you could please stay on the page
for at least 30 seconds to help me support
my work as per Medium's new changes.
```

Scenario 1: Question: You are responsible for optimizing costs in an AWS environment. How would you go about identifying and reducing unnecessary expenses?

Answer: To optimize costs in AWS, I would start by utilizing AWS Cost Explorer to analyze the cost and usage data. I would identify underutilized or idle resources and either modify, resize, or terminate them. Additionally, I would implement AWS Budgets and set up alerts to receive notifications when costs exceed predefined

thresholds. Implementing auto-scaling for resources based on demand can also help optimize costs by automatically adjusting capacity.

Scenario 2: Question: Your company is planning to migrate its on-premises database to Amazon RDS. Outline the key steps you would take to ensure a successful database migration.

Answer: The key steps for a successful database migration to Amazon RDS include:

1. Assessment: Understand the existing database schema, dependencies, and performance metrics.
2. Schema Conversion: Modify the schema to be compatible with the target RDS engine.
3. Data Migration: Use AWS Database Migration Service (DMS) to migrate data with minimal downtime.
4. Testing: Conduct thorough testing to ensure data integrity, performance, and functionality.
5. Cutover: Plan and execute the cutover to switch production traffic to the RDS instance.
6. Post-Migration Validation: Perform post-migration checks to ensure everything is functioning as expected.

Scenario 3: Question: Your application, hosted on AWS, is experiencing a sudden increase in traffic. How would you handle this surge in demand and ensure that the application remains responsive and available?

Answer: To handle a sudden increase in traffic, I would:

1. Auto-scaling: Configure auto-scaling groups to automatically adjust the number of EC2 instances based on traffic.
2. Content Delivery Network (CDN): Utilize AWS CloudFront to cache and deliver content closer to end-users, reducing the load on the origin server.
3. Load Balancing: Implement an Elastic Load Balancer (ELB) to distribute incoming traffic across multiple instances.
4. Caching: Implement caching mechanisms, such as Amazon ElastiCache, to reduce the load on the database and improve response times.

5. Monitoring: Set up CloudWatch alarms to receive notifications and automatically trigger scaling actions based on predefined thresholds.

Scenario 4: Question: Your team is working on a microservices architecture, and you need to implement communication between microservices. How would you approach this using AWS services?

Answer: I would implement communication between microservices using the following AWS services:

1. Amazon API Gateway: Use API Gateway to create RESTful APIs that expose microservices to clients.
2. AWS Lambda: Deploy serverless functions for business logic, allowing for scalable and cost-effective execution.
3. Amazon SNS (Simple Notification Service): Implement SNS for event-driven communication between microservices.
4. Amazon SQS (Simple Queue Service): Use SQS to decouple and manage messages between microservices.
5. Amazon EventBridge: Utilize EventBridge for event-driven architectures, enabling communication through events and rules.

Scenario 5: Question: Your application requires a highly available and fault-tolerant storage solution. How would you design a scalable storage architecture using AWS services?

Answer: For a highly available and fault-tolerant storage architecture, I would:

1. Amazon S3 (Simple Storage Service): Use S3 for durable, scalable, and low-latency object storage.
2. Cross-Region Replication: Implement cross-region replication for redundancy and disaster recovery.
3. Versioning: Enable versioning in S3 to track and recover previous versions of objects.
4. Amazon Glacier: Utilize Glacier for long-term archival of infrequently accessed data.
5. Amazon EFS (Elastic File System): Implement EFS for scalable and shared file storage across multiple instances.

Scenario 6: Question: Your team is planning to deploy a highly available web application on AWS. How would you design the architecture to ensure resilience and minimize downtime?

Answer: To design a highly available web application on AWS, I would:

number of EC2 instances based on demand, ensuring availability.

3. Multi-AZ Deployments: Deploy resources across multiple Availability Zones to ensure resilience against failures in a single zone.
4. Elastic Load Balancer (ELB): Use ELB to distribute incoming traffic across multiple instances, improving fault tolerance.
5. Amazon RDS Multi-AZ: Deploy the database using Multi-AZ configuration for automatic failover and increased availability.

Scenario 7: Question: Your organization is looking to enhance security by implementing encryption for data at rest. How would you implement encryption in AWS?

Answer: To implement encryption for data at rest in AWS:

1. Amazon S3 Server-Side Encryption: Enable server-side encryption for S3 buckets using AWS Key Management Service (KMS) or Amazon S3-managed keys (SSE-S3).
2. Amazon EBS Volume Encryption: Encrypt Amazon EBS volumes attached to EC2 instances using AWS KMS.
3. Amazon RDS Encryption: Enable encryption for RDS databases using AWS KMS for enhanced security.
4. AWS KMS Customer Master Keys: Manage and rotate Customer Master Keys (CMKs) in AWS KMS regularly for improved key security.
5. AWS CloudHSM: For additional security, consider using AWS CloudHSM for dedicated hardware security module (HSM) protection.

Scenario 8: Question: Your application processes sensitive data, and compliance is a top priority. How would you ensure that your AWS environment meets compliance requirements?

Answer: To ensure compliance in an AWS environment:

1. AWS Artifact: Leverage AWS Artifact to access compliance reports and documentation to support audit requirements.
2. AWS Config: Implement AWS Config to assess, audit, and evaluate the configurations of AWS resources against predefined rules.
3. AWS CloudTrail: Enable CloudTrail to log and monitor all API calls, providing visibility into user activity.
4. AWS Identity and Access Management (IAM): Define and enforce strict IAM policies to control access to resources based on the principle of least privilege.
5. Encryption and Security Best Practices: Adhere to encryption standards, implement network security, and follow other security best practices outlined in AWS documentation.

Scenario 9: Question: Your team is considering the adoption of a serverless architecture for specific workloads. What are the advantages and challenges of using AWS Lambda?

Answer: Advantages:

1. Scalability: AWS Lambda automatically scales based on demand, allowing for efficient resource utilization.
2. Cost Savings: Pay only for the compute time consumed, leading to cost-effective solutions for intermittent workloads.
3. Event-Driven: Easily integrate with various AWS services and trigger functions based on events.
4. Managed Service: AWS Lambda is a fully managed service, eliminating the need for server provisioning and maintenance.

Challenges:

1. Cold Start Latency: There may be initial latency (cold start) when a function is invoked after being idle.
2. Execution Time Limits: Functions are limited by maximum execution times, and long-running processes may require a different approach.
3. Limited State: Designed for stateless functions; maintaining state requires additional considerations.

4. Debugging Complexity: Debugging serverless applications can be challenging compared to traditional architectures.

Scenario 10: Question: Your team is planning to deploy a containerized application on AWS. How would you manage container orchestration and scaling?

Answer:

1. Amazon ECS (Elastic Container Service): Use ECS to deploy, manage, and scale containers with ease. ECS provides a fully managed container orchestration service.
2. Amazon EKS (Elastic Kubernetes Service): Consider EKS for Kubernetes-based container orchestration, offering scalability and flexibility.
3. Auto Scaling Groups: Implement Auto Scaling Groups to automatically adjust the number of EC2 instances in the container cluster based on demand.
4. Amazon EC2 Spot Instances: Use Spot Instances to reduce costs for non-critical workloads, taking advantage of spare EC2 capacity.
5. Application Load Balancer: Deploy an Application Load Balancer to distribute incoming traffic across containers, ensuring optimal load balancing.

> ***P*lease consider giving it a clap and follow to help me reach more like-minded individuals. Thank you for being a part of this community!**

Interview Jobs AWS DevOps Sre

Follow

Written by Deep

1K Followers

Cloud / DevOps Engineer | I write on Cloud, AWS, K8s, DevOps and CI/CD Buy-Me-A-Coffee
https://buymeacoffee.com/deepdevops

More from Deep

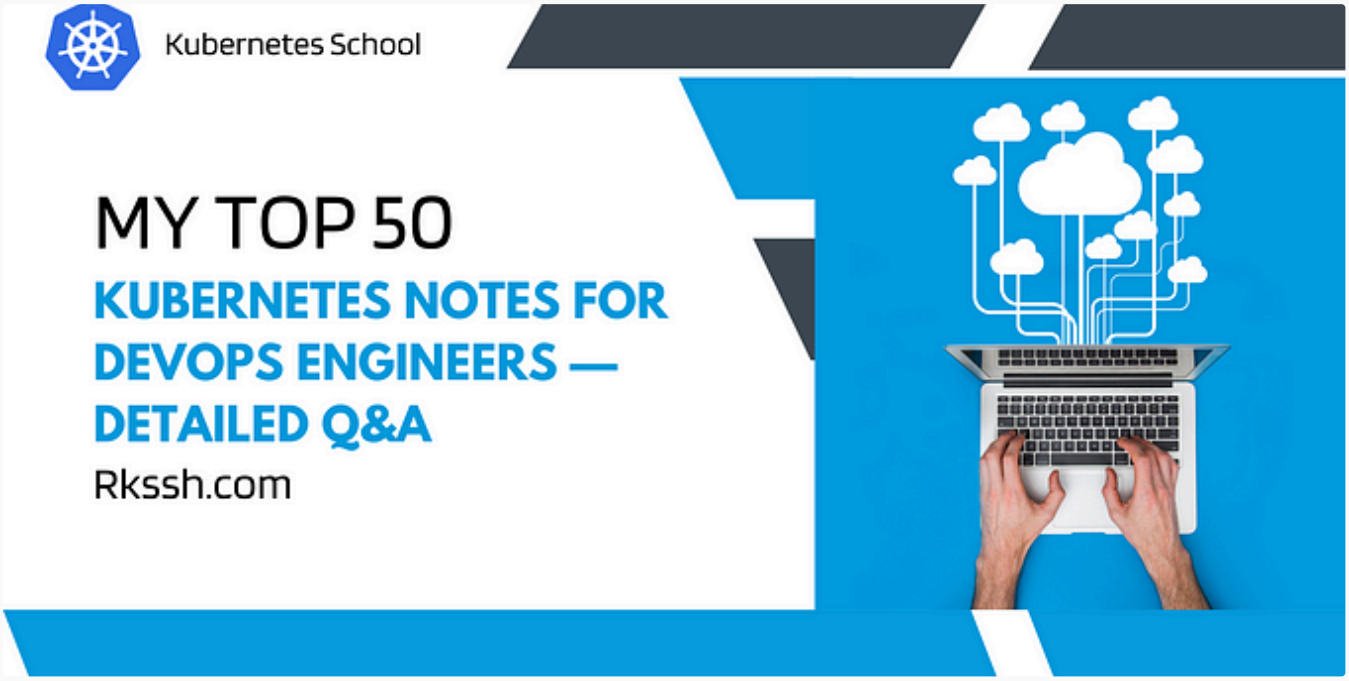

Deep in CodeX

My Top 50 Kubernetes Notes for DevOps Engineers—Detailed Q&A

My Notes Series Kubernetes

Aug 9 477 2

Deep

My Top 50 Terraform Notes for DevOps Engineers—Detailed Q&A

On a free Medium plan? Click here

Aug 28 162

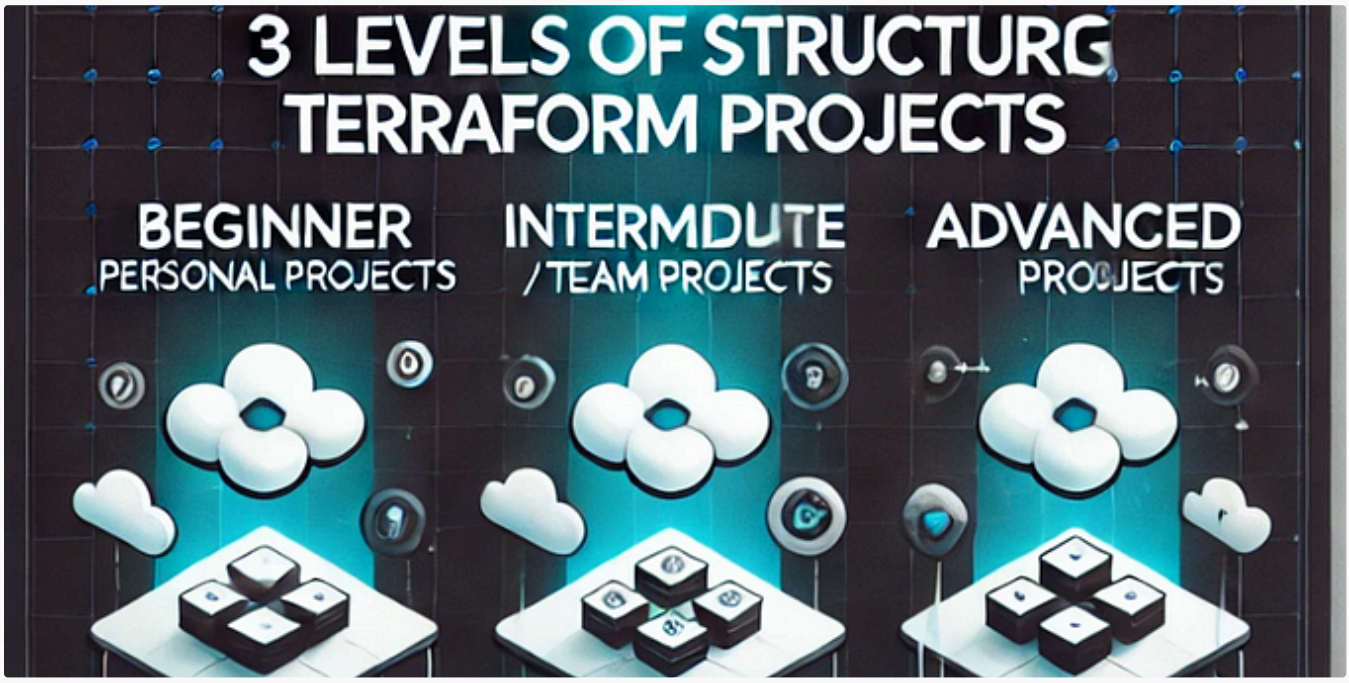

Deep in DevOps.dev

3 Levels of structure terraform projects

The Evolution of a Terraform Project - From Tiny Side Project to Big-Time Enterprise

Aug 13 70 4

Deep in CodeX

The Toughest Kubernetes Interview Questions I Faced — You Won't Believe

From managing clusters to scaling nightmares, here are the trickiest Kubernetes questions that left me stumped during my interviews.

Sep 4 22 1

See all from Deep

Recommended from Medium

Nurunnubi Talukder in Cloud, DevOps, Security & AI Career Talk

Solution Architect AWS interview questions and answers[25Q-Part1]!!

Aspiring to become a Solution Architect on AWS? This guide covers essential interview questions and answers to help you prepare. Whether...

Jun 30 43

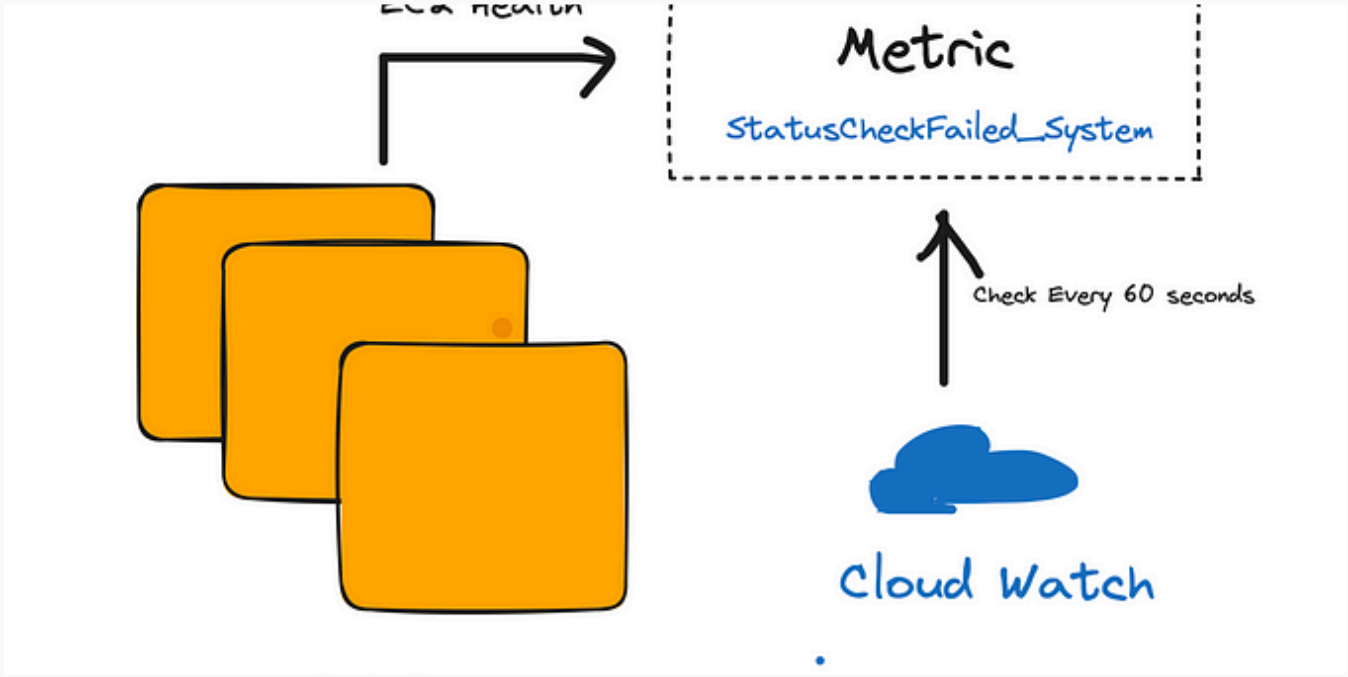

Vikas Taank

Orchestrating Cloud Watch Alarm for EC2

AWS does recover failed virtual machines under only some circumstances. For example, AWS will not recover an EC2 instance if a whole rack...

Jun 21 1

Lists

Stories to Help You Level-Up at Work

19 stories · 793 saves

The New Chatbots: ChatGPT, Bard, and Beyond

12 stories · 456 saves

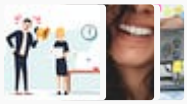

Work 101

26 stories · 167 saves

Stories to Help You Grow as a Software Developer

19 stories · 1349 saves

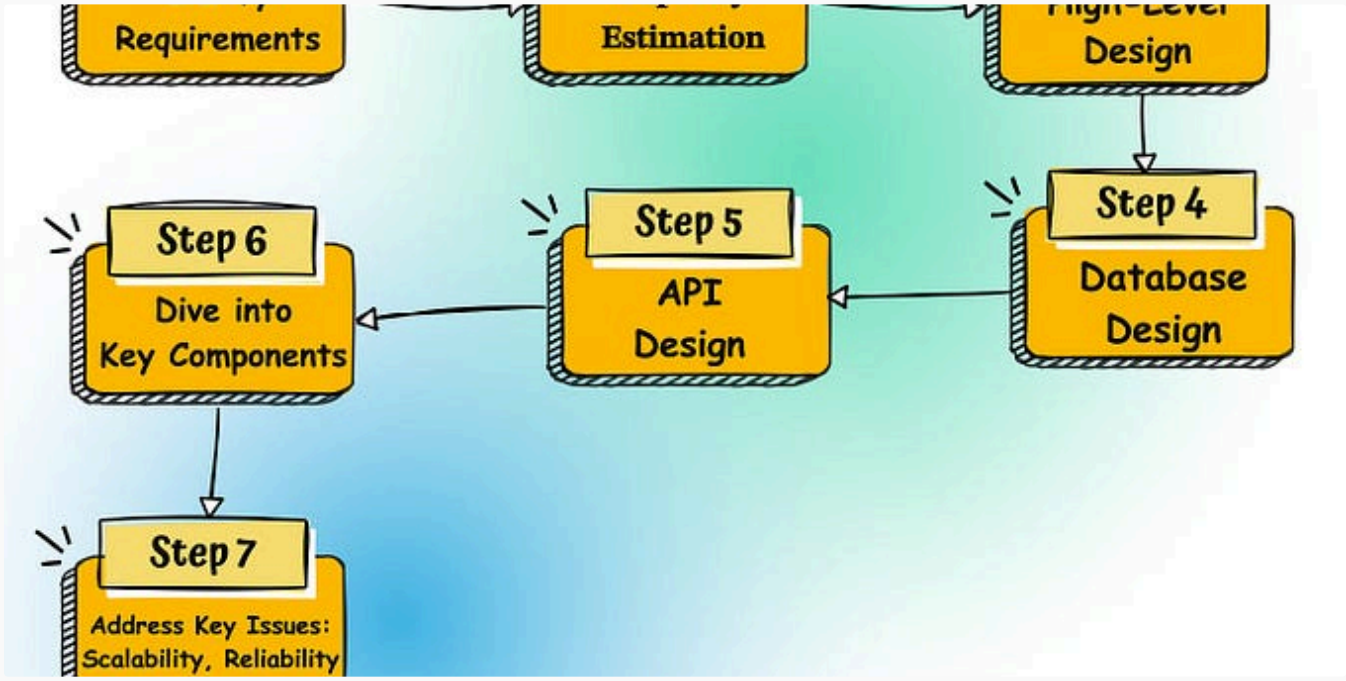

Ashish Pratap Singh in AlgoMaster.io

A Step-by-Step Guide to System Design Interviews

System design interviews can be one of the most challenging but with a structured approach, you can navigate these interviews more...

Aug 25 123 1

Ashwin

AWS EC2 Instance Storage [Interview Series]

Are you familiar with EC2 instance storage, EBS volumes, and Amazon EFS? If not, don't worry!

Apr 29 1

Liu Zuo Lin in Level Up Coding

AWS Data Engineer Associate Cheatsheet (Consolidated) + Self-Test

If you're studying for the AWS Data Engineer Associate Certification DEA-C01, hopefully this is helpful in some way. I'm also assuming that...

Apr 2 720 1

Aditya Pathak

Docker Interview Questions Basic to Advanced

Summary about docker.

Mar 14 1